

AWS Lambda Overview

Run code without provisioning or managing servers

Introduction to AWS Lambda

- Function-as-a-Service
- Run code without provisioning or managing servers
- Pay only for the compute time you consume
- Automatically runs your code with high availability
- Scale with usage

Lambda handles

- Load balancing
 - Auto scaling
 - Handling failures
 - Security isolation
 - OS management
 - Managing utilization
- (and many other things) for you

What is Serverless?




No infrastructure provisioning,
no management



Automatic scaling

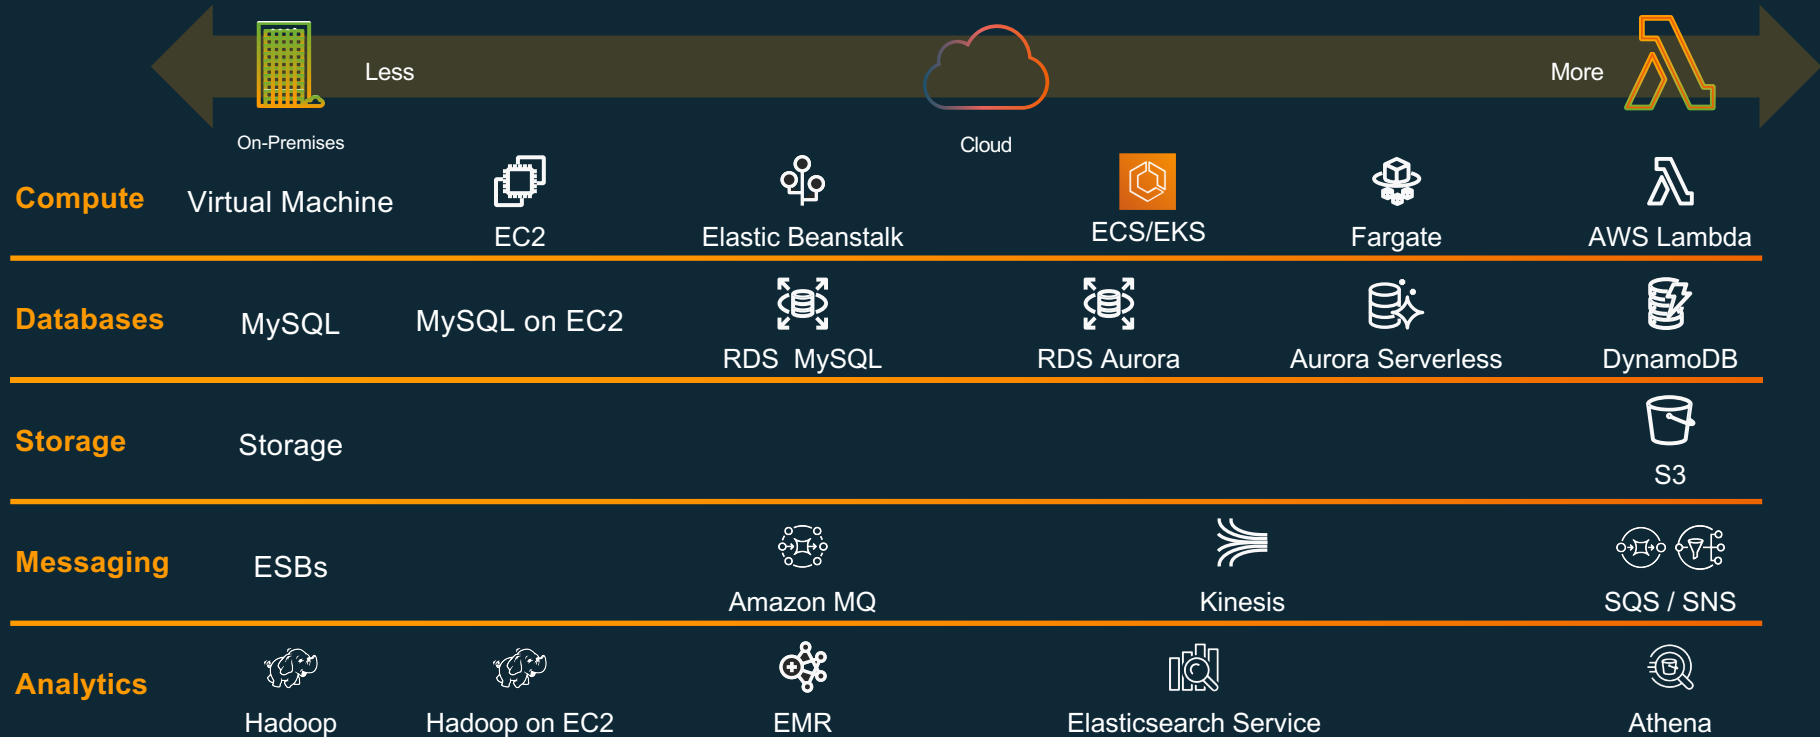
Pay for value



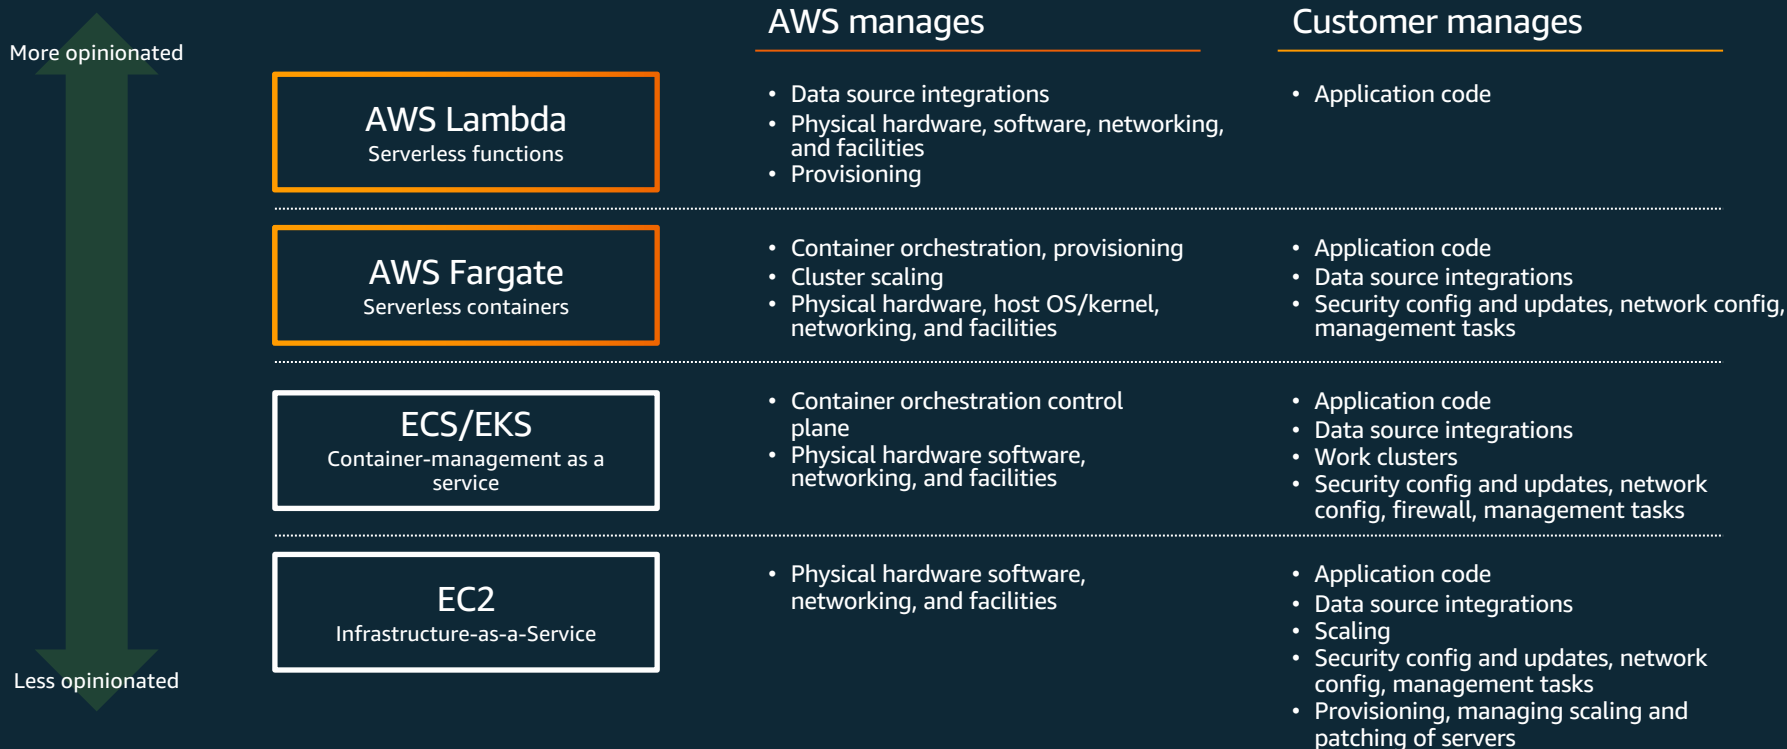
Highly available and secure



AWS operational responsibility models



Comparison of operational responsibility



Anatomy of a Lambda Function

Serverless applications



Serverless applications

Function



Node.js
Python
Java
C#
Go
Ruby
Runtime API

Serverless Applications

Event source



Function



Changes in
data state



Requests to
endpoints

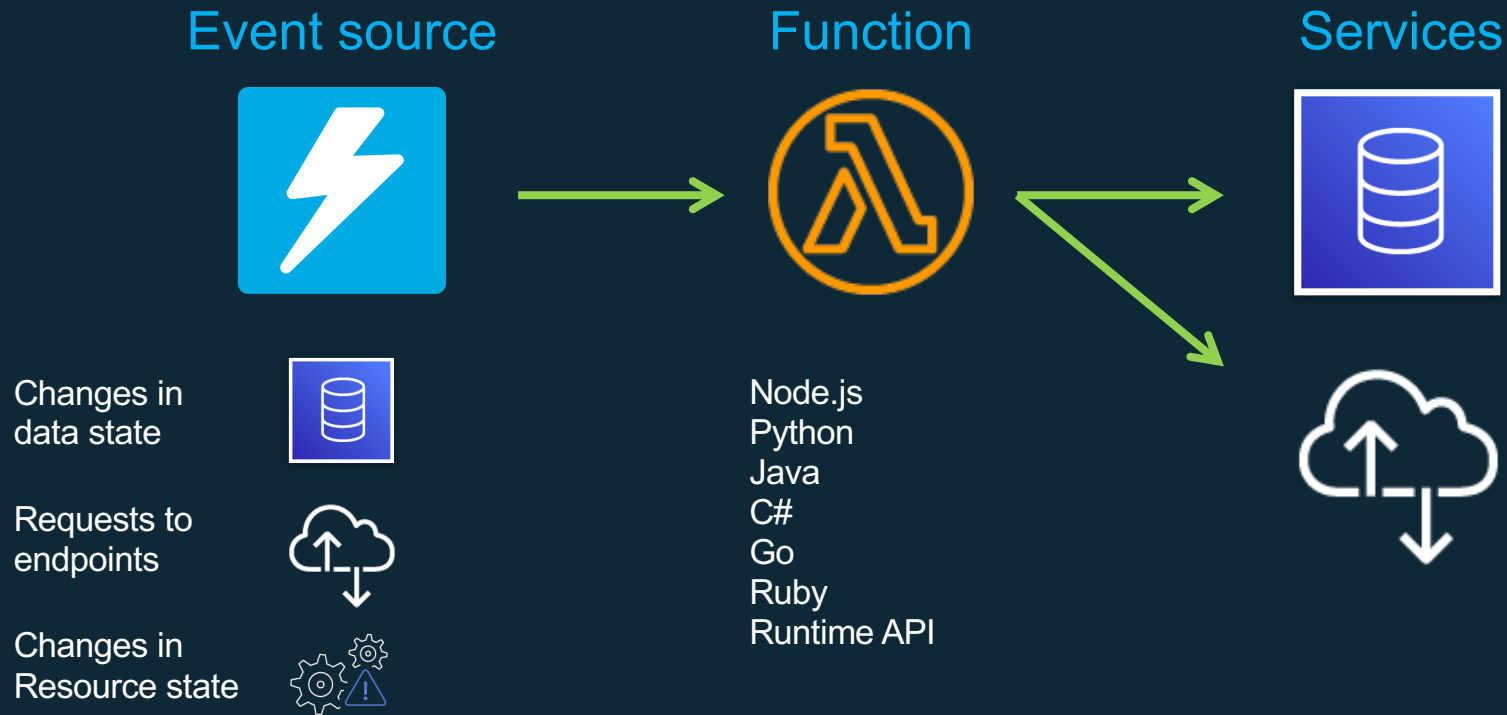


Changes in
Resource state



Node.js
Python
Java
C#
Go
Ruby
Runtime API

Serverless Applications



Anatomy of a Lambda Function

Handler() function

Function to be executed upon invocation

Event object

Data sent during Lambda function Invocation

Context object

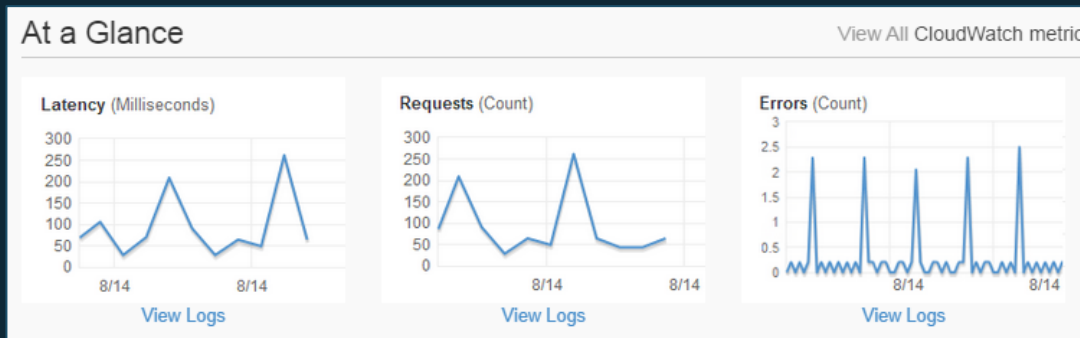
Methods available to interact with runtime information (request ID, log group, more)

```
import json

def lambda_handler(event, context):
    # TODO implement
    return {
        'statusCode': 200,
        'body': json.dumps('Hello world!')
    }
```

Monitoring and debugging Lambda functions

- AWS Lambda console includes a dashboard for functions
 - Lists all Lambda functions
 - Easy editing of resources, event sources and other settings
 - At-a-glance metrics
- Metrics automatically reported to Amazon CloudWatch for each Lambda function
 - Requests
 - Errors
 - Latency
 - Throttles



Lambda Layers



Lets functions easily share code: Upload layer once, reference within any function

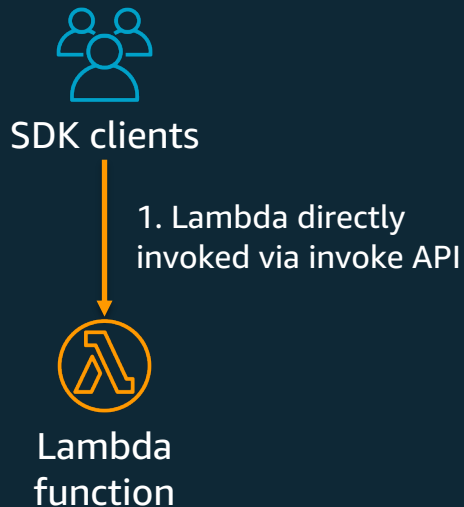
Promote separation of responsibilities, lets developers iterate faster on writing business logic

Built in support for secure sharing by ecosystem

Lambda Layers: Uses cases

- Custom code, that is used by more than one function
- Libraries, modules, frameworks to simplify the implementation of your business logic
 - Security/monitoring service
- Shared code that does not change frequently
- Bring your own Runtime
 - C++
 - Rust
 - PHP

Lambda API



API provided by the Lambda service

Used by all other services that invoke Lambda across all models


Supports sync and async calls

Can pass any event payload structure you want


Client included in every SDK

Lambda Execution Models


Synchronous (push)



Asynchronous (event)



Stream (Poll-based)



Dead-Letter Queue

- Asynchronous Lambda invocations are retried two more times (3 times total)
- Lambda can forward payloads that were not processed to a dead-letter queue (IF configured!)
- A mechanism to handle exceptions and failures gracefully


Lambda Permissions Model

Function policies:


- “Actions on bucket X can invoke Lambda function Z”
- Resource policies allow for cross account access
- Used for sync and async invocations

Execution role:

- “Lambda function A can read from DynamoDB table users”
- Define what AWS resources/API calls can this function access via IAM
- Used in streaming invocations



Fine-grained Pricing



Free Tier

1M requests and 400,000 GBs of compute.
Every month, every customer.

Buy compute time in 100ms increments

Low request charge

No hourly, daily, or monthly minimums

No per-device fees


Never pay for idle

Example event sources that trigger AWS Lambda

DATA STORES




Amazon S3



Amazon
DynamoDB



Amazon
Kinesis



Amazon
Cognito

ENDPOINTS



Amazon
API Gateway



AWS IoT



AWS Step
Functions




Amazon
Alexa

CONFIGURATION REPOSITORIES




AWS
CloudFormation



AWS CloudTrail




AWS
CodeCommit




Amazon
CloudWatch

EVENT/MESSAGE SERVICES



Amazon
SES



Amazon SNS



Cron events

... and more on the way!

Common AWS Lambda use cases



Web Apps

- Static websites
- Complex web apps
- Packages for Flask and Express




Backends

- Apps & services
- Mobile
- IoT



Data Processing

- Real time
- MapReduce
- Batch



Chatbots

- Powering chatbot logic



Amazon Alexa

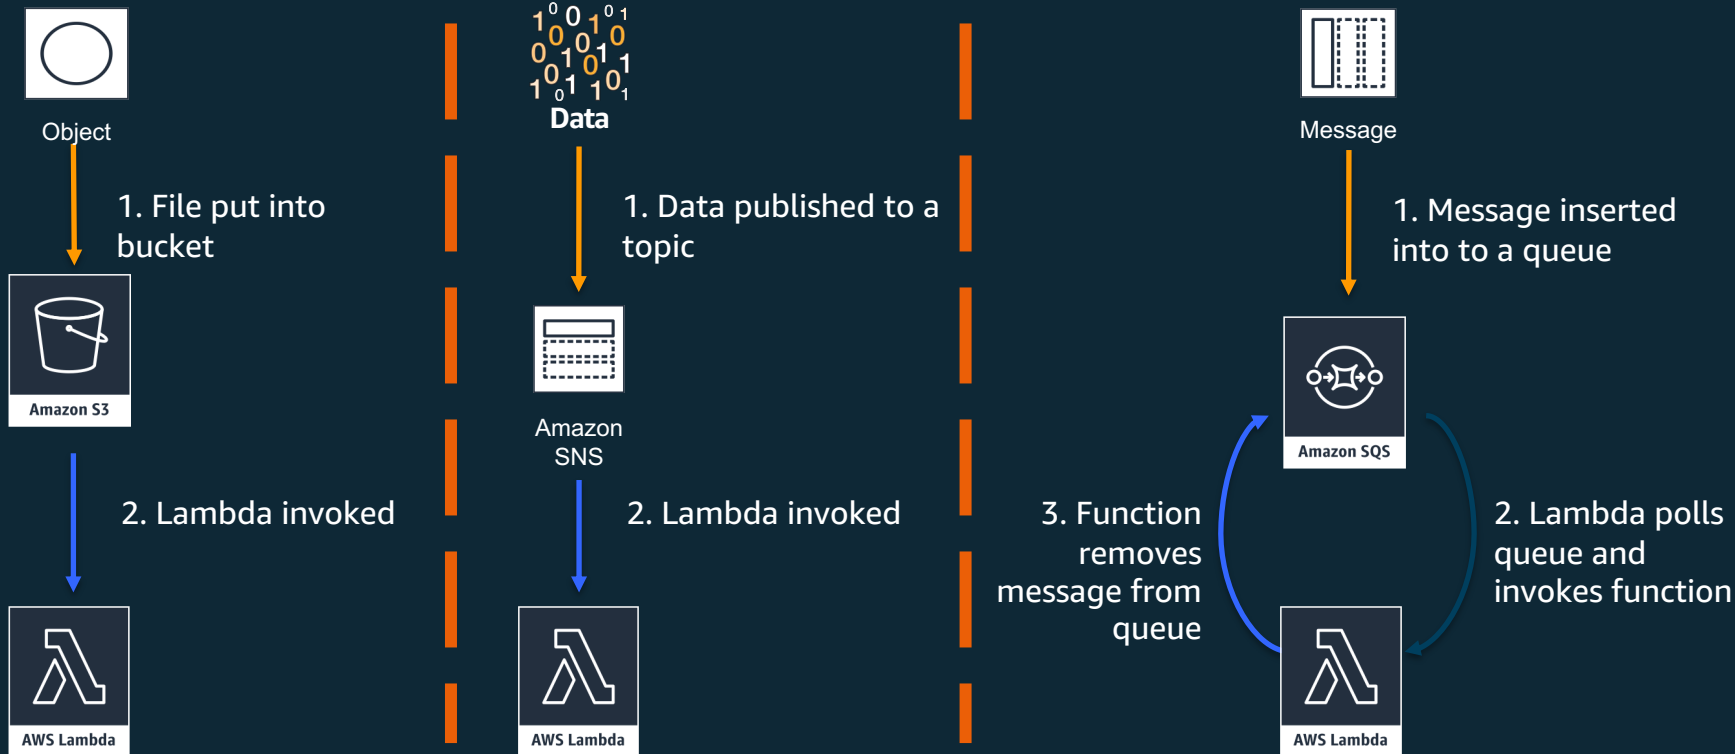
- Powering voice-enabled apps
- Alexa Skills Kit



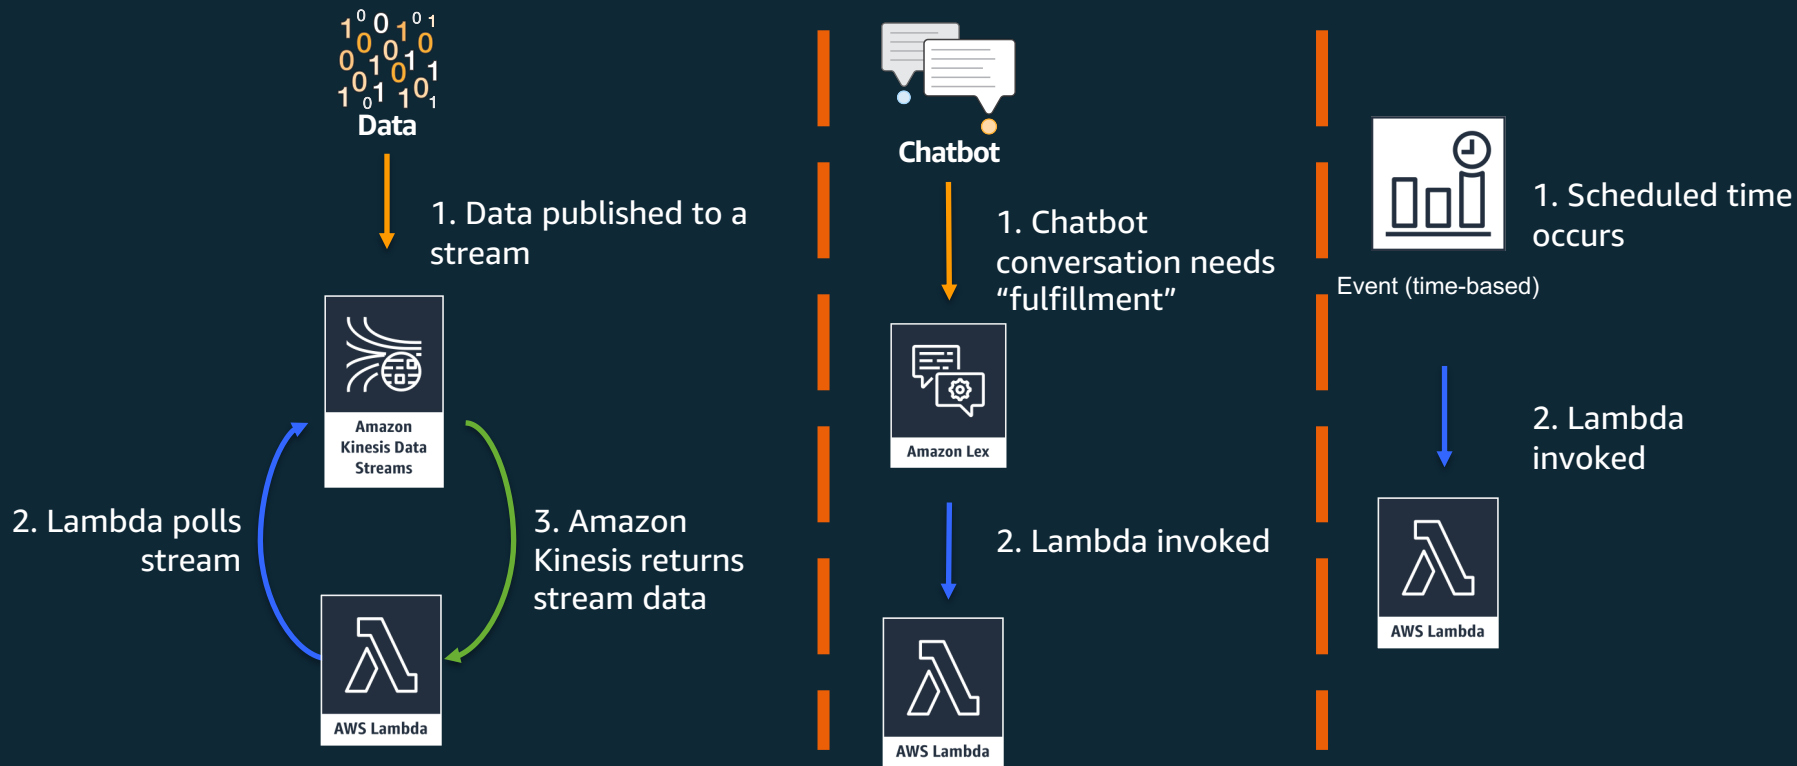
IT Automation

- Policy engines
- Extending AWS services
- Infrastructure management

Serverless architectures



Serverless Architectures



When do we choose Lambda over other compute offerings?

AWS Compute Offerings



Service

Amazon EC2

AWS Fargate

AWS Lambda

Unit of scale

VM

Task

Function

Level of
abstraction

H/W

OS

Runtime

AWS Compute Offerings



Service

Amazon EC2

How do I
choose?

I want to
configure
servers, storage,
networking, and
my OS



AWS Fargate

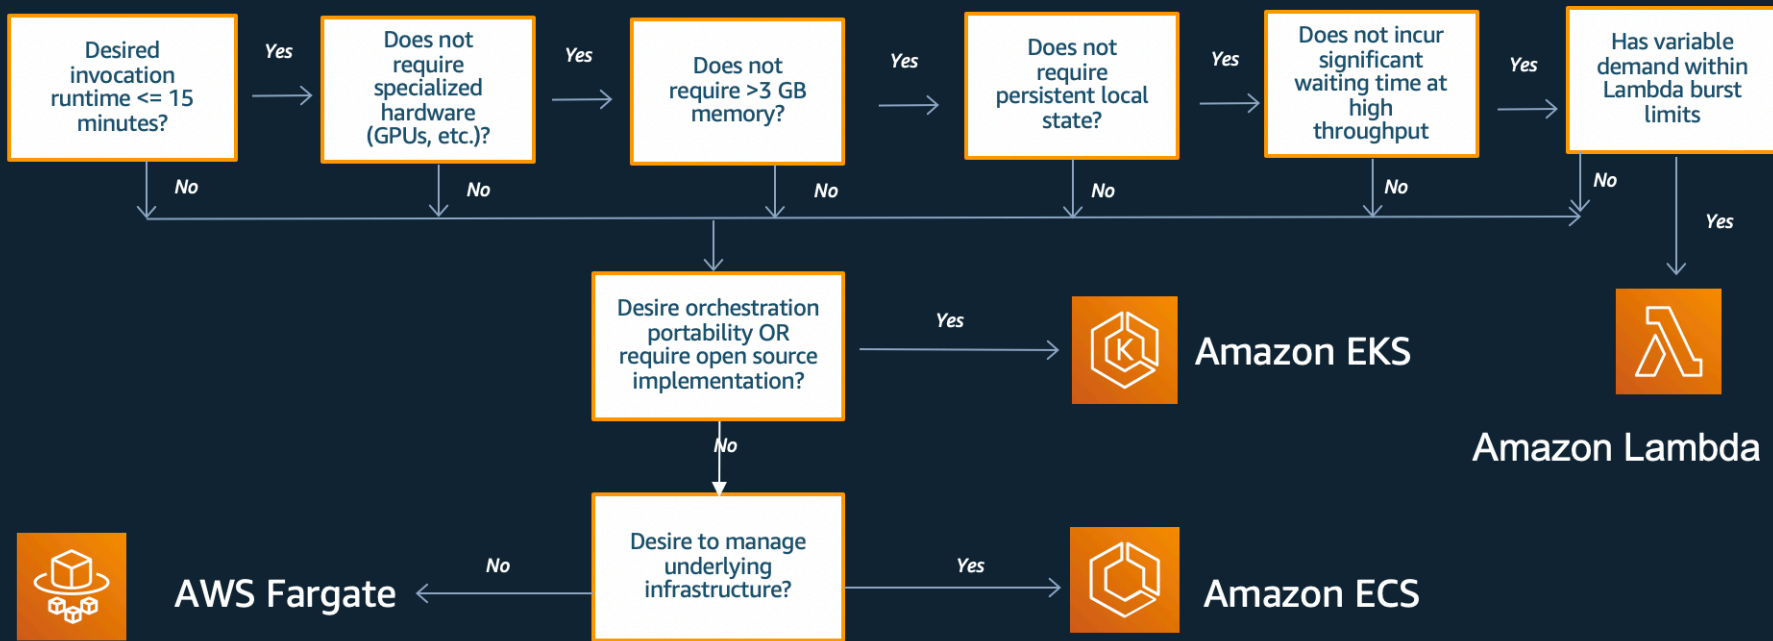
I want to run my
containers



AWS Lambda

Run my code
when it's
needed

Picking a Compute Platform: Containers vs. Lambda



Step Functions

Orchestration for serverless apps

“I want to sequence functions”

“I want to select functions based on data”

“I want to run functions in parallel”

“I want to retry functions”

“I want to try/catch/finally”

“I want to run code for hours”



AWS Step Functions

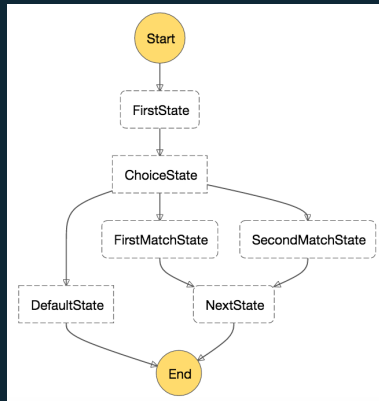
AWS Step Functions

Easily coordinate multiple Lambda functions using visual workflows

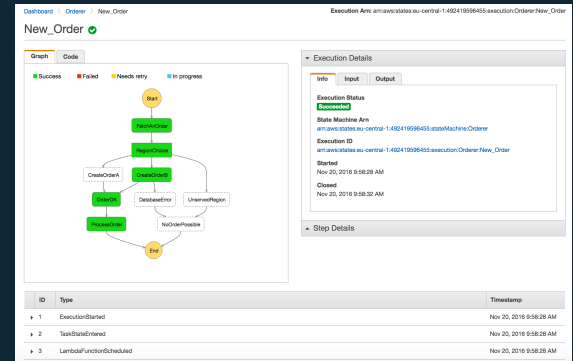
Define in JSON

```
Code
1 {
2   "Comment": "An AWS example using a choice state.",
3   "StartAt": "FirstState",
4   "States": {
5     "FirstState": {
6       "Type": "Task",
7       "Resource": "arn:aws:lambda:REGION:ACCOUNT_ID:function:FUNCTION_NAME",
8       "Next": "ChoiceState"
9     },
10    "ChoiceState": {
11      "Type": "Choice",
12      "Choices": [
13        {
14          "Variable": "$?.Resource",
15          "Match": "Equals",
16          "Next": "FirstMatchState"
17        },
18        {
19          "Variable": "$?.Resource",
20          "Match": "NotEquals",
21          "Next": "SecondMatchState"
22        },
23        {
24          "Default": "DefaultState"
25        }
26      ]
27    },
28    "FirstMatchState": {
29      "Type": "Task",
30      "Resource": "arn:aws:lambda:REGION:ACCOUNT_ID:function:FUNCTION_NAME",
31      "Next": "NextState"
32    },
33    "SecondMatchState": {
34      "Type": "Task",
35      "Resource": "arn:aws:lambda:REGION:ACCOUNT_ID:function:FUNCTION_NAME",
36      "Next": "NextState"
37    },
38    "DefaultState": {
39      "Type": "Task",
40      "Resource": "arn:aws:lambda:REGION:ACCOUNT_ID:function:FUNCTION_NAME",
41      "Next": "NextState"
42    },
43    "NextState": {
44      "Type": "Task",
45      "Resource": "arn:aws:lambda:REGION:ACCOUNT_ID:function:FUNCTION_NAME",
46      "Next": "End"
47    },
48    "End": {
49      "Type": "End"
50    }
51  }
52 }
```

Visualize in the
Console



Monitor
Executions



Benefits of Step Functions orchestration

Productivity




Coordinate and visualize Lambda functions as a series of steps to quickly create serverless apps

Resilience



Automatically trigger and track each step at scale and handle errors with built-in retry and fallback

Agility



Change and add steps without writing code to evolve applications and innovate faster


AWS Service integrations with Step Functions



AWS Lambda



AWS Batch




Amazon DynamoDB



AWS Fargate



Amazon Elastic
Container Service



Amazon Simple
Notification Service




AWS Step Functions



Amazon Simple
Queue Service



Amazon SageMaker



AWS Glue

Serverless Application Model

AWS Serverless Application Model (AWS SAM)

AWS CloudFormation brings:

- Infrastructure as code
- Easy to provision and manage a collection of related AWS resources
- Input .yaml file and output provisioned AWS resources
- Optimized for infrastructure

AWS SAM:

- CloudFormation extension optimized for serverless
- New serverless resources: functions, APIs, and tables
- Supports anything CloudFormation supports
- Open specification (Apache 2.0)



AWS Lambda best practices

- Limit your function/code size
- Node – remember execution is asynchronous
- 500 MB /tmp directory provided to each function
- Don't assume function will reuse underlying infrastructure
 - But take advantage of it when it does occur
- You own the logs
 - Include details from service-provided context
- Create custom metrics
 - Operations-centric vs. business-centric

Additional best practices

- Use environment variables
 - Parameterize code and change parameters independent of code updates
 - Use for securing credentials and keeping them out of code
- Externalize authorization to IAM roles whenever possible
 - Least privilege and separate IAM roles
- Externalize configuration
 - DynamoDB is great for this
- Take advantage of dead letter queues
 - Use to handle failed invocations

Additional best practices

- Make sure your downstream setup “keeps up” with Lambda scaling
 - Limit concurrency when talking to relational databases
- Be aware of service throttling
 - Engage AWS Support to increase your limits
- Contact AWS Support before known large scaling event
 - Infrastructure Event Management (IEM) offers real-time support for large scaling events
 - IEM is available for Enterprise and Business support customers

Thank you!