



# Solutions Architect Professional

AWS Scaling and Resiliency



# Cloud Architecture Principles

## Goal as an architecture

Deliver Highly Available,  
Highly Scalable,  
Fault Tolerant,  
High Performance,  
Secure Cloud Architecture


# Cloud Architecture Principles

“Everything fails, all the time.”

*Werner Vogels, CTO, Amazon.com*

# Cloud Architecture Principles


## 1. Design for failure



Goal: Applications should continue to function even if the underlying application component fails, communication is lost or physical hardware fails, is removed/replaced.

# Cloud Architecture Principles

## 2. Embrace Elasticity & Automate



- Do not assume health, availability or fixed location of components
- Automate installation and configuration of environment
- Favor dynamic configuration


# Cloud Architecture Principles

## 3. Loosely Couple



# Cloud Architecture Principles


## 3. Loosely Couple



- Design architectures with independent components
- Design every component as a black box
- Load balance clusters

# Cloud Architecture Principles

## 3. Loosely Couple




- Use queues to pass messages between components



# Cloud Architecture Principles


## 4. Become Stateless



- Don't store state in server
- Leverage services to hold state information
- Application functions regardless of which application node processes the request

# Cloud Architecture Principles


## 5. Leverage Parallelism



- One Server working for Four hours costs the same as Four servers working for One hour
- Combine with elasticity to increase capacity when you need it most

# Cloud Architecture Principles

## 6. Use appropriate storage options



- Don't log clicks to RDBMS, use NoSQL data store
- Don't store images in RDBMS, use object store
- Offload log files to scalable object storage

# Cloud Architecture Principles

## 7. Build Security into every layer



- Encrypt data in transit and rest between application tiers
- Enforce principle of least privilege across every service
- Automatically rotate security keys frequently

# From Development to Web-Scale

Development

Test

Production

Web-Scale

Beyond

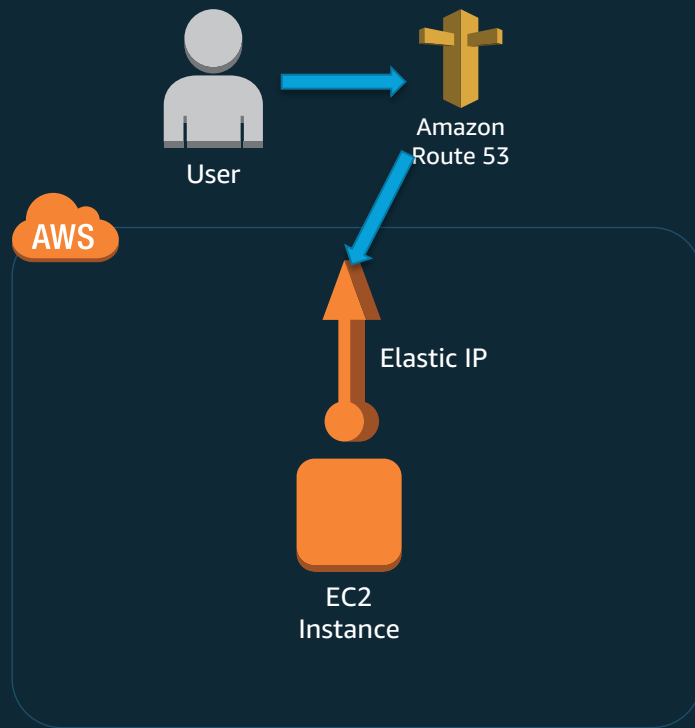


# Development

Your own development environment:

- An EC2 Instance, hosting
  - Web, App
  - Database
  - Management
  - Etc.

A single Elastic IP  
Route53 for DNS



# "We're gonna need a bigger box"


## Different EC2 instance type

- High memory instances
- High CPU instances
- High I/O instances
- High storage instances

Can now leverage PIOPs

Easy to change instance sizes

Will hit an endpoint eventually



# "We're gonna need a bigger box"


## Different EC2 instance type

- High memory instances
- High CPU instances
- High I/O instances
- High storage instances

Can now leverage PIOPs

Easy to change instance sizes

Will hit an endpoint eventually






# Shall we use this to test with?

There are some issues


- Constrained by a single environment
- Too many eggs in one basket
- No failover
- No redundancy



# Shall we use this to test with?

There are some issues

- Constrained by a single environment
- Too many eggs in one basket
- No failover
- No redundancy



# From Development to Web-Scale

Development

Test

Production

Web-Scale

Beyond



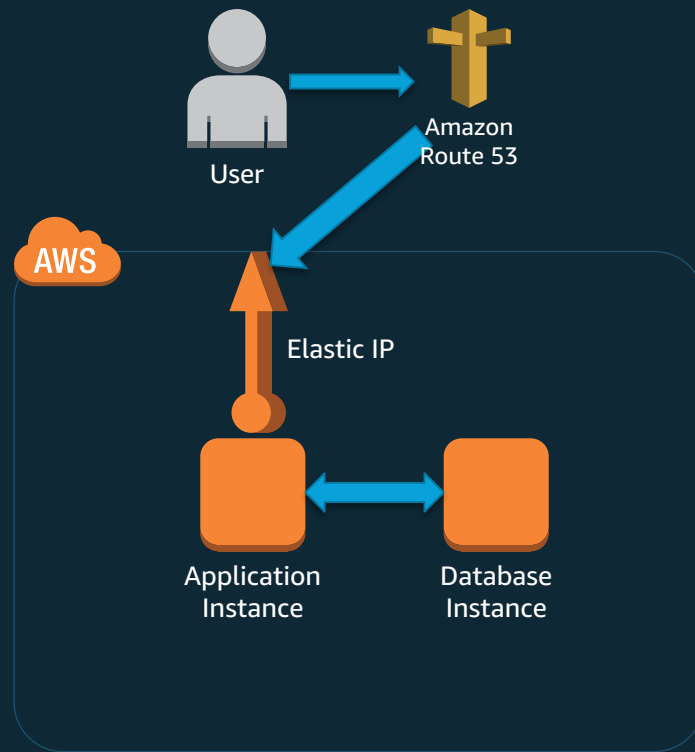
# Test

First let's separate out our single host into different tiers:

Web

Database

Should we make use of a database service?



# AWS Database Options

## Self-managed



### Database Server on Amazon EC2

Your choice of database running on Amazon EC2

Bring Your Own License (BYOL)

## Fully Managed



### Amazon RDS

Microsoft SQL, Oracle PostgreSQL or MySQL as a managed service

Flexible licensing – BYOL or license included



### Amazon DynamoDB

Managed NoSQL database service using SSD storage

Seamless scalability  
Zero administration



### Amazon Redshift

Massively parallel, petabyte-scale, data warehouse service

Fast, powerful and easy to scale

# Why start with SQL?

Established and well worn technology

Lots of existing code, communities, books, background, tools, etc

You aren't going to break SQL DBs until you're really big

- But you might break parts of it (hence blended approach)

Clear patterns to scalability

- But, can suffer vertical scaling constraints ultimately

# Why else might you need NoSQL?

Super low latency applications

Metadata driven datasets

Highly un-relational data

Need schema-less data constructs\*

Massive amounts of data (again, in the TB range)

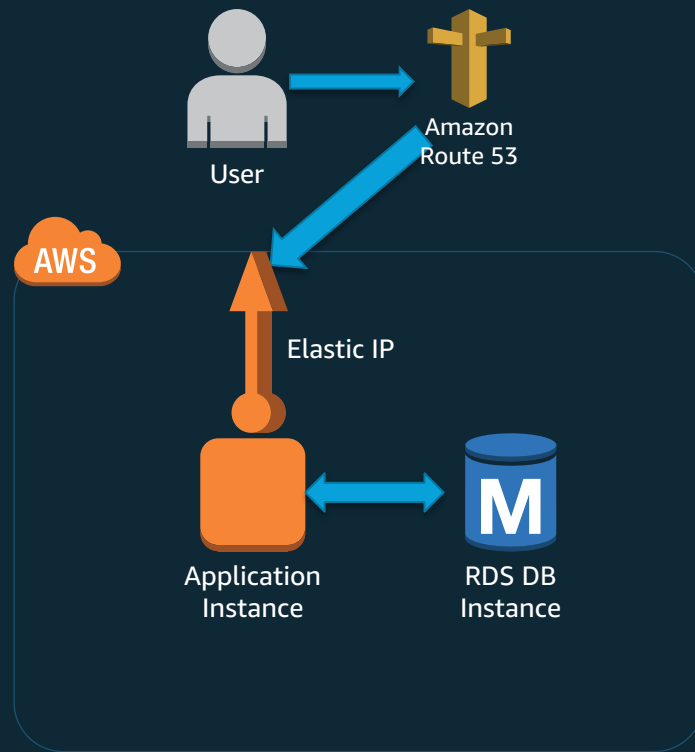
Rapid ingest of data (thousands of records/sec)

\*Need != "its easier to do dev without schemas"

# Test

Lets leverage RDS for our database tier minimize the infrastructure management

Now our test environment is available as a separate set of infrastructure





# From Development to Web-Scale

Development

Test

Production

Web-Scale

Beyond



# Production V1

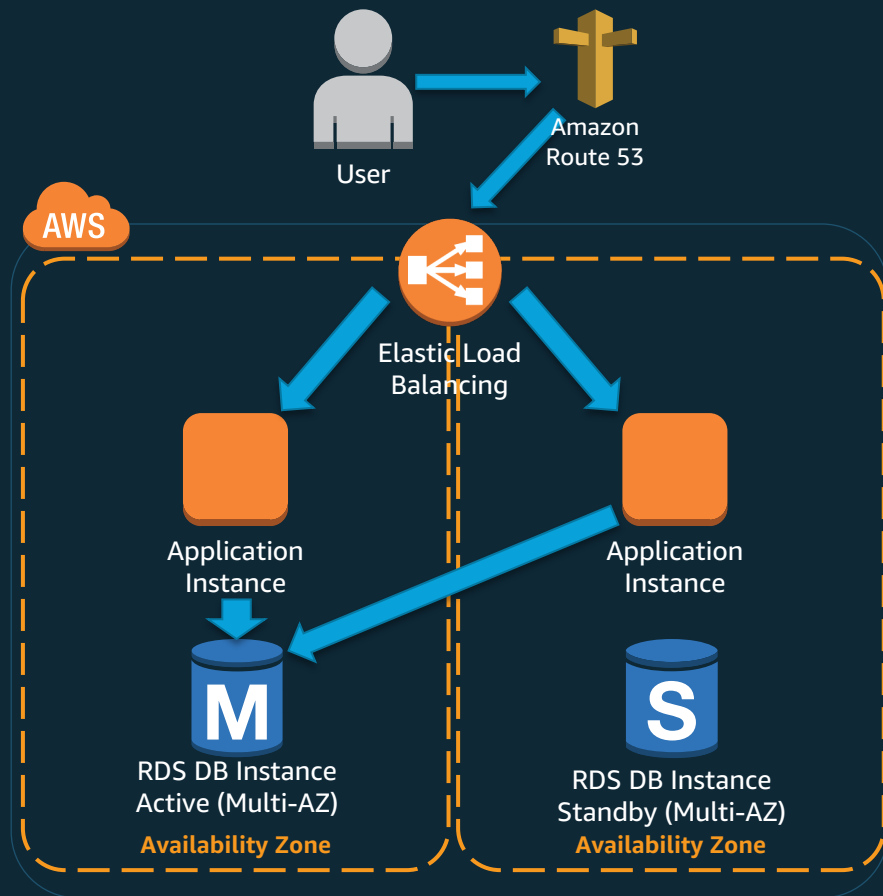
Now we're ready to go to production, we need to address our lack of failover and redundancy issues

Another Application Instance


- In another Availability Zone

Elastic Load Balancing

Enable Amazon RDS multi-AZ



# Elastic Load Balancing



- Load Balancing as a service
- Automatically distributes incoming application traffic across multiple Amazon EC2 instances
- Enables you to achieve greater levels of fault tolerance in your applications
- Built-in application health detection, serve traffic only to operational application instances
- Provides SSL offload
- Seamlessly provides required amount of capacity needed to distribute application traffic


# RDS Availability & Scaling Options

Multi-AZ Replication across availability zones of master DB

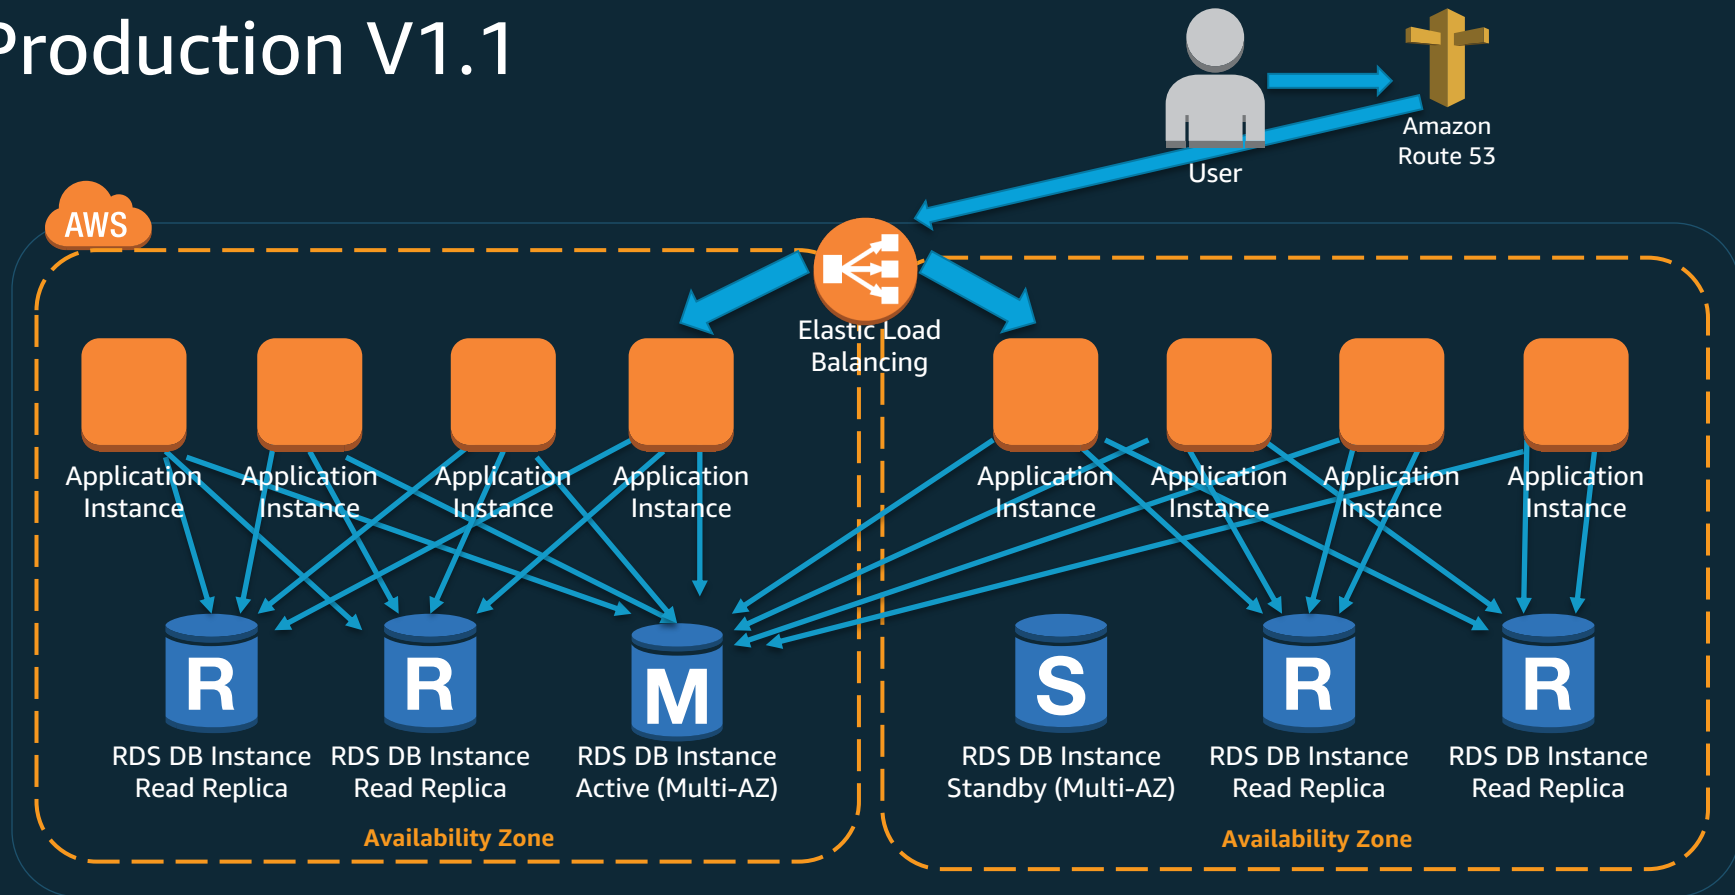
Automated failover from master to slave

Vertically scale RDS by changing instance size

Horizontally scale RDS with Read-Replicas



# Production V1.1



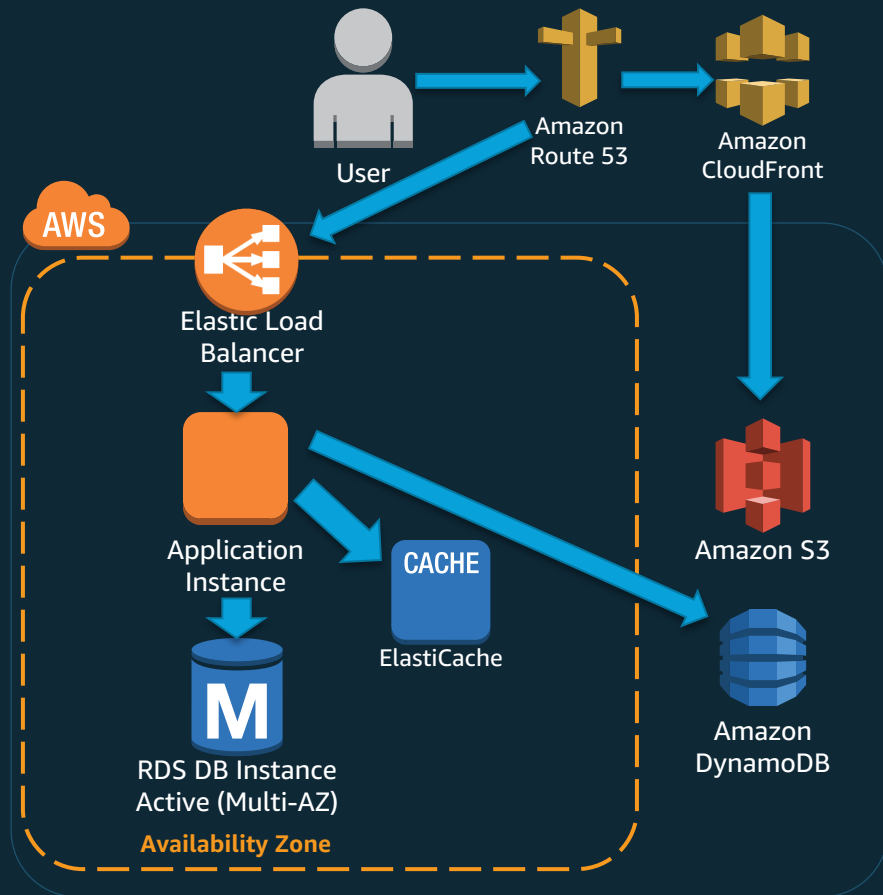
# Shift Some Load Around

Let's lighten the load on our web and database instances:

Move static content from the Application Instance to Amazon S3 and CloudFront

Move dynamic content from the Elastic Load Balancing to CloudFront

Move session/state and DB caching to ElastiCache or DynamoDB



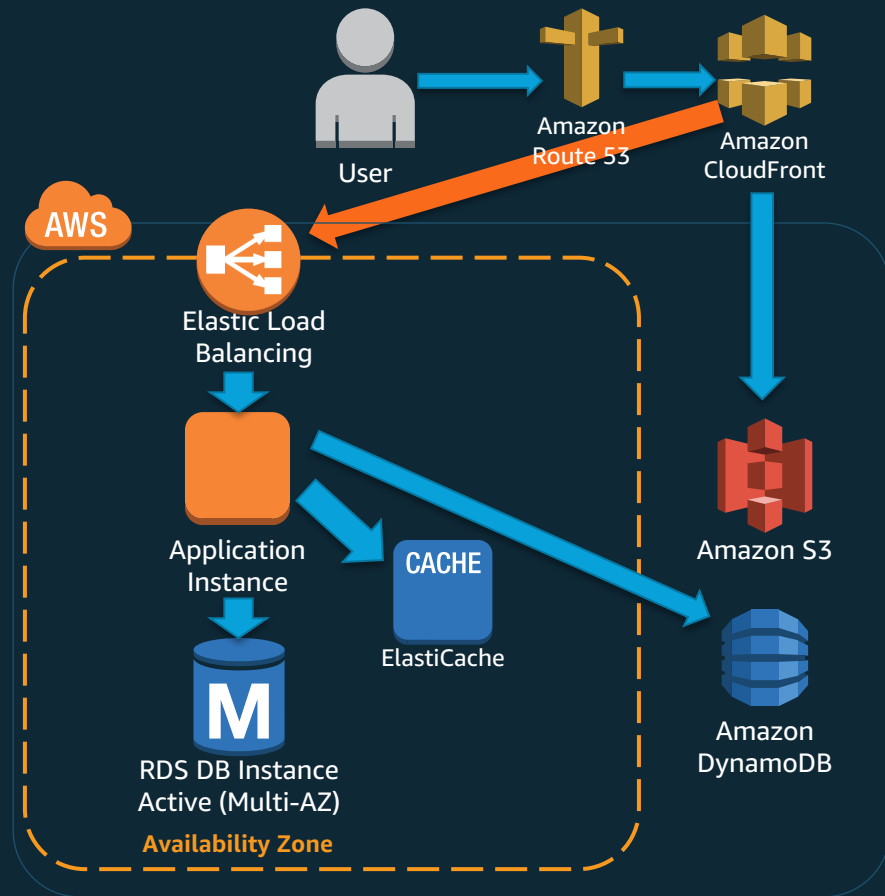
# Shift Some Load Around

Let's lighten the load on our web and database instances:

Move static content from the Application Instance to Amazon S3 and CloudFront

Move dynamic content from the Elastic Load Balancing to CloudFront

Move session/state and DB caching to ElastiCache or DynamoDB



# Amazon CloudFront

Amazon CloudFront is a web service for scalable content delivery.

Cache content at the edge locations around the world for faster delivery

Helps lower load on origin infrastructure

Dynamic and static content


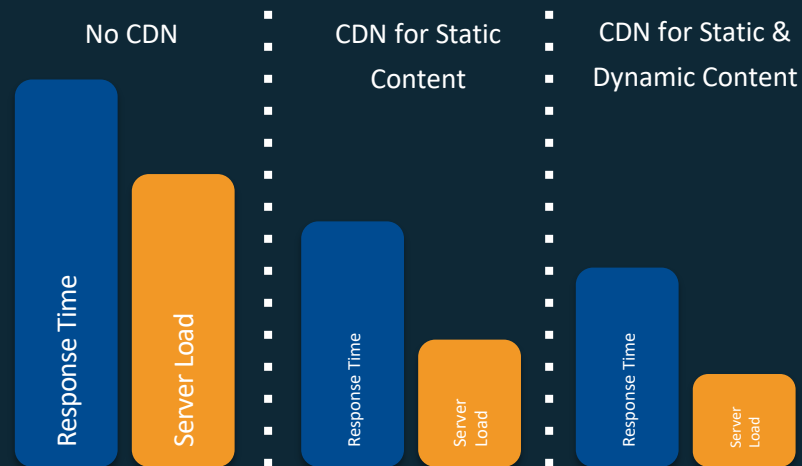
Streaming video

Zone apex support

Custom SSL certificates

Low TTLs ( as short as 0 seconds )

Optimized to work with EC2, Amazon S3, Elastic Load Balancing, and Route53





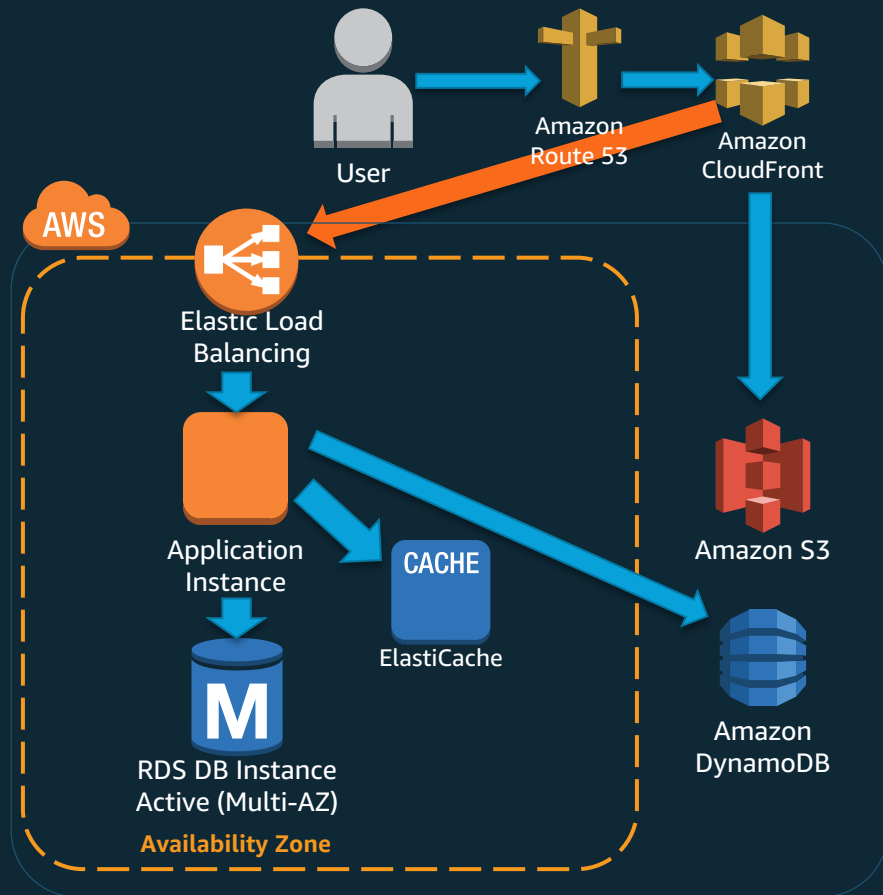
# Shift Some Load Around

Let's lighten the load on our web and database instances

Move static content from the Application Instance to Amazon S3 and CloudFront

Move dynamic content from the Elastic Load Balancing to CloudFront

Move session/state and DB caching to ElastiCache or Amazon DynamoDB



# ElastiCache

## Hosted Memcached & Redis

- Speaks same API as traditional open source Memcached and Redis

Scale from one to many nodes

Self-healing ( replaces dead instance )

Very fast ( single digit ms speeds usually )

Local to a single AZ for Memcache, with no persistence or replication

With Redis can put a replica in a different AZ with persistence

Use AWS's Auto Discovery client to simplify clusters growing and shrinking, without affecting your application



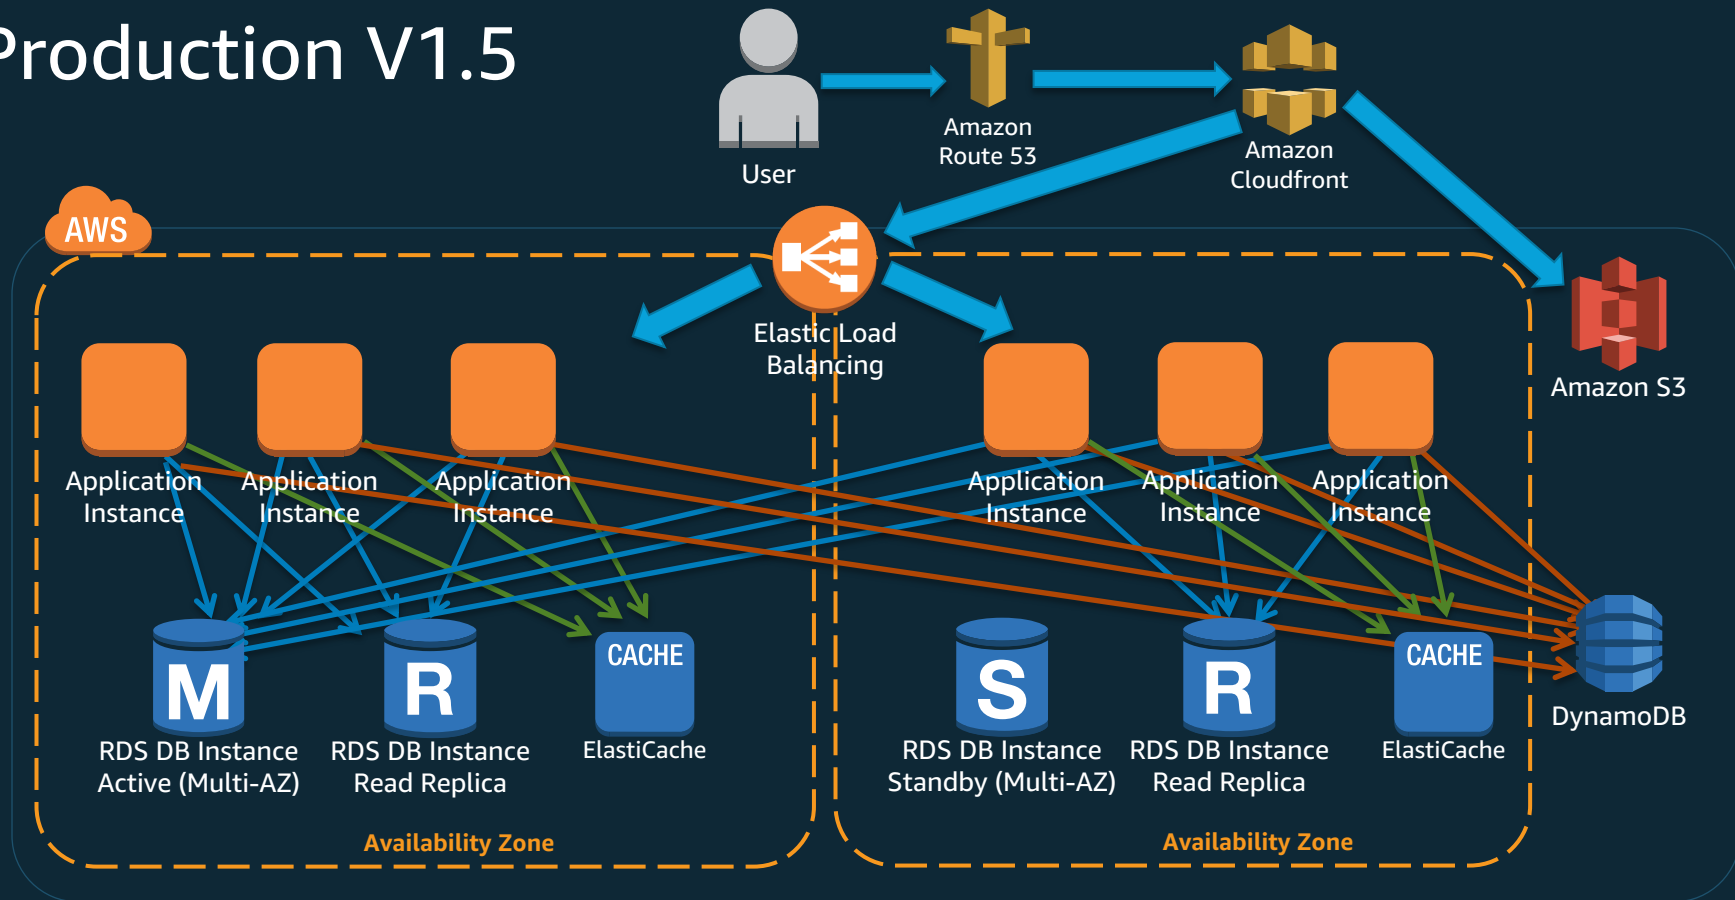
# Amazon DynamoDB

- Fully managed, provisioned throughput NoSQL database
- Fast, predictable performance
- Fully distributed, auto-partitioning, fault-tolerant architecture
- Consider for non-uniform data



Feature	Details
<b>Provisioned throughput</b>	Dial up or down provisioned read/write capacity
<b>Predictable performance</b>	Average single-digit millisecond latencies from SSD-backed infrastructure
<b>Strong consistency</b>	Be sure you are reading the most up to date values
<b>Fault tolerant</b>	Data replicated across Availability Zones
<b>Monitoring</b>	Integrated to CloudWatch
<b>Secure</b>	Integrates with AWS Identity and Access Management (IAM)

# Production V1.5



# Elasticity with Auto-Scaling

# Auto Scaling


Automatic resizing of compute clusters based on demand

Define minimum and maximum number of instances

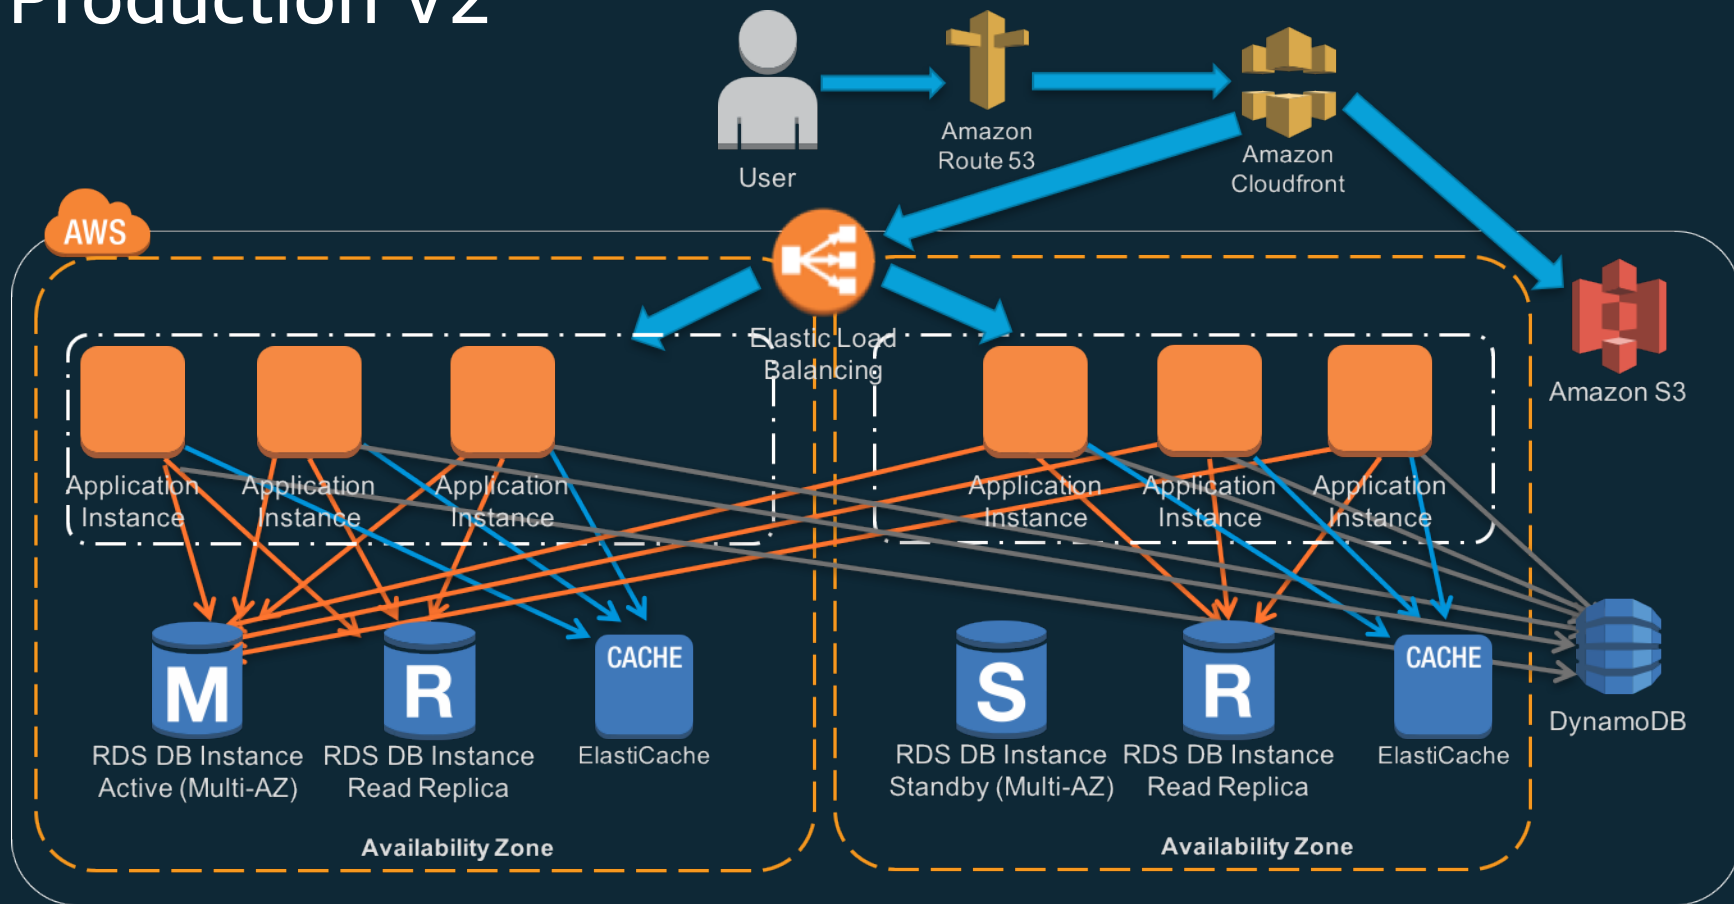
Define when scaling up and down occurs

Use metrics collected in Amazon CloudWatch to drive scaling

Run Auto Scaling for On-Demand and Spot instance types




# Production V2



# AWS Application Management Solutions

Higher level services

Do it yourself



Elastic Beanstalk



AWS OpsWorks



AWS CloudFormation



EC2

Convenience

Control



# From Development to Web-Scale

Development

Test

Production

Web-Scale

Beyond



# Service Oriented Architecture

# SOA

## Service Oriented Architecture

- Move services into their own tiers or modules
- Treat each of these as 100% separate pieces of your architecture
- Scale them independently

Amazon.com and AWS do this extensively! It offers flexibility and greater understanding of each component.

# Cloud SOA

## Service Oriented Architecture

- Same SOA principles, but leveraging Cloud services
- Allows you to architect for
  - Automation
  - Scale
  - Cost
  - Availability
- Applications are
  - Always-on
  - Self-healing
- Cloud-native New Applications

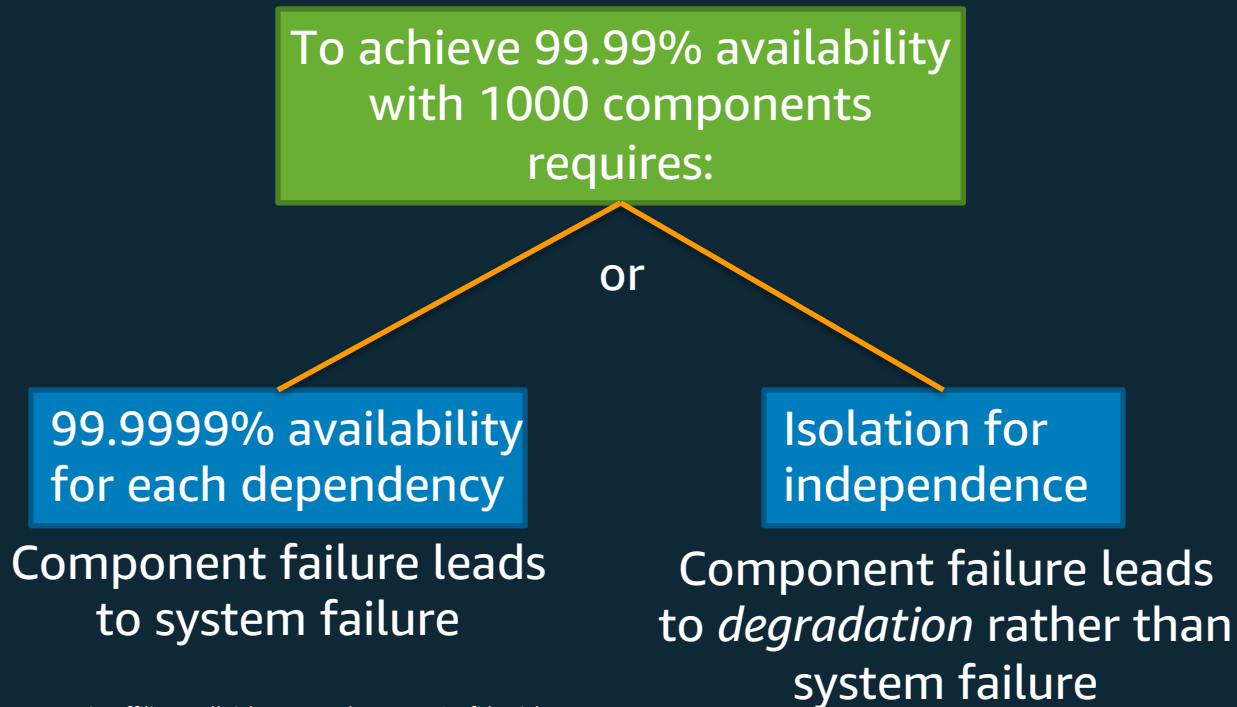
# Availability Conundrum



<i>N</i> Service Dependencies	Availability
2	99.98%
10	99.9%
100	99%
1000	90%

*N* dependencies

# Availabilities Compound



# Loose Coupling

The looser they're coupled, the bigger they scale

- Use independent components
- Design everything as a black box
- Decouple interactions
- Favor services with built in redundancy and scalability than building your own

## Use Amazon SQS as Buffers

Tight Coupling



Loose Coupling



# Loose Coupling + Cloud SOA = Winning

Examples:

Email

Queuing

Transcoding

Search

Databases

Monitoring

Metrics

Logging



Amazon SNS



Amazon  
ElasticSearch



Amazon SQS




Amazon SES



Amazon SWF

If someone has a service for it already, use that instead of building it yourself





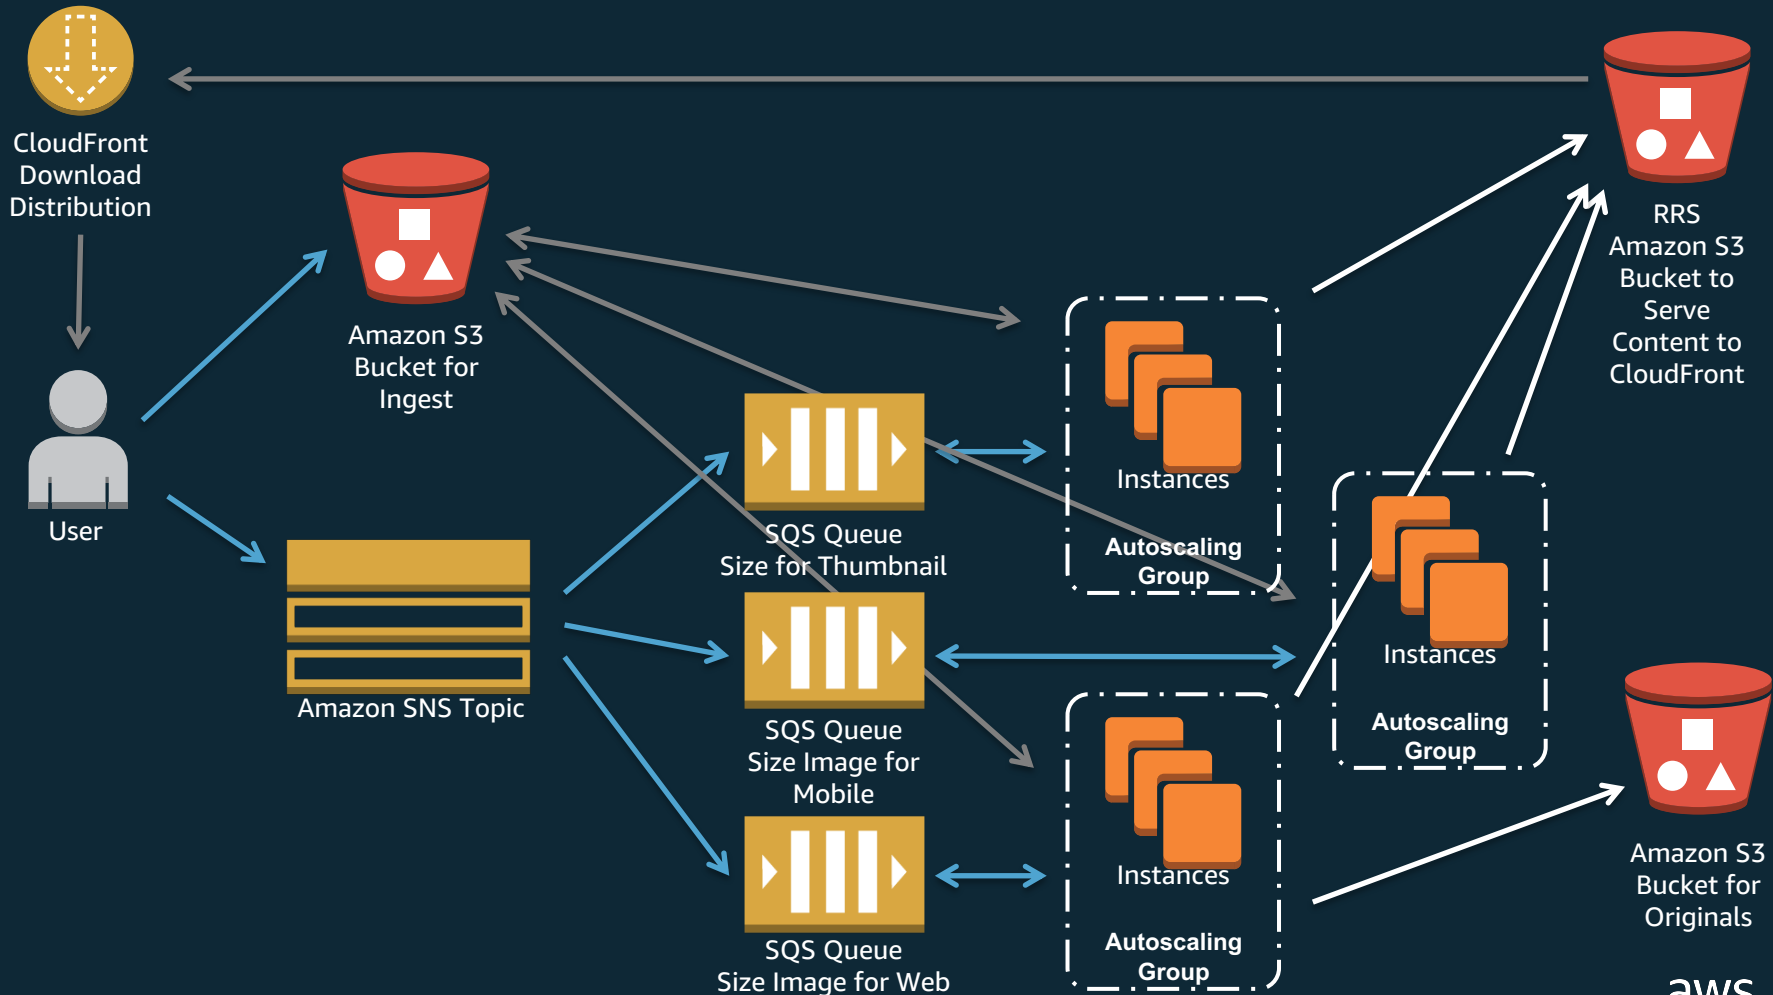


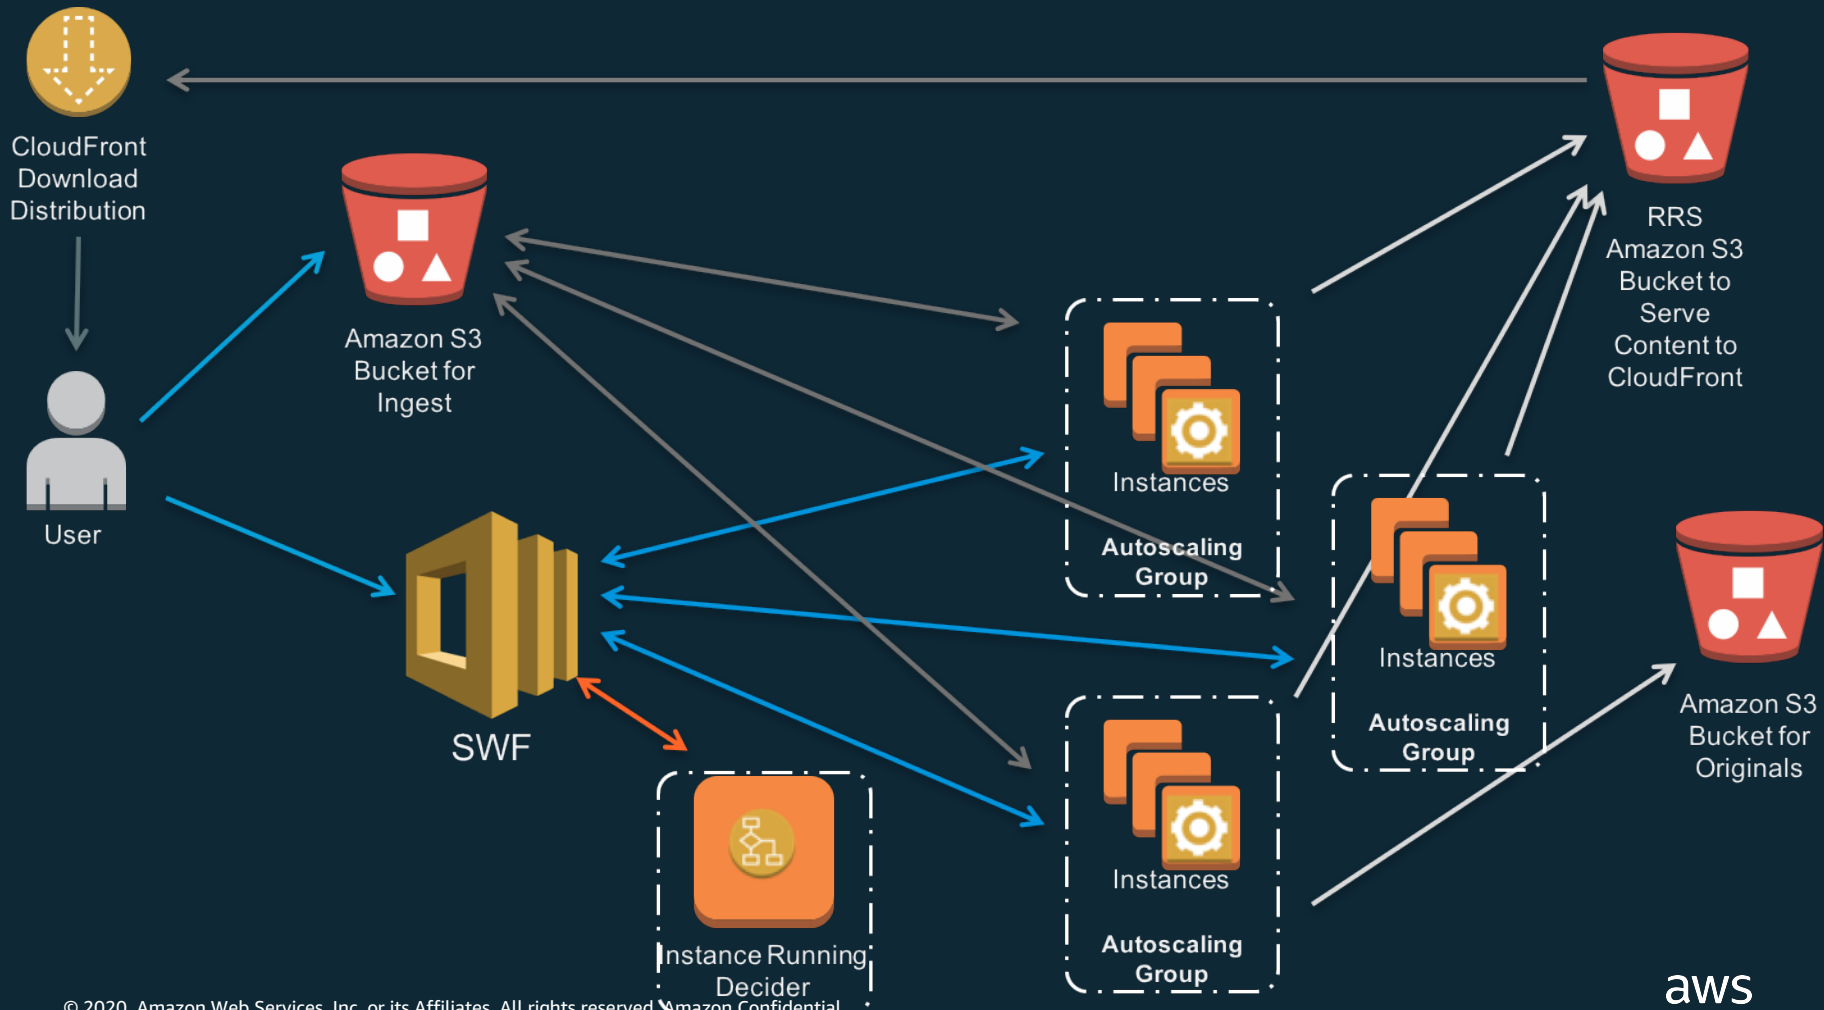
# Amazon Simple Workflow Service



Provides an orchestration tool across your infrastructure  
Can act as a middle layer to pass messages and setup tasks  
Lets you break down individual tasks into different workers  
Lets you define logic between workers  
Lets you make a worker task from anything that can be scripted  
Includes built-in retries, timeouts, logging  
Features built-in reliability, scalability, and low cost

Your code =  &   
Deciders Workers





# Web-Scale

Reaching web-scale will require most or all of the following:

- Multi-AZ
- Elastic Load Balancing between tiers
- Auto Scaling
- Service-oriented architecture
- Serving content smartly (S3/CloudFront)
- Caching database queries
- Moving state off tiers that auto-scale

# Web-Scale

